

Week 3 Tasks Solutions

```
library(tidyverse)
library(tidymodels)
library(moderndiver)
library(gapminder)
library(skimr)
library(mvtnorm)
library(gridExtra)
library(ISLR)
```

Tasks


1. Examine the relationship between teaching score and age in the `evals` data set. What is the value of the correlation coefficient? How would you interpret this verbally? Finally, produce a scatterplot of teaching score and age.

```
evals.age <- evals |>
  select(score, age)
evals.age |>
  get_correlation(formula = score ~ age)
```

```
# A tibble: 1 x 1
  cor
  <dbl>
1 -0.107
```

```
ggplot(evals.age, aes(x = age, y = score)) +
  geom_point() +
  labs(x = "Age", y = "Teaching Score",
```

```
title = "Relationship between Teaching Score and Age")
```



2. Perform a formal analysis of the relationship between teaching score and age by fitting a simple linear regression model. Superimpose your best-fitting line onto your scatterplot from Task 1.

```
evals.age <- evals |>
  select(score, age)
model <- linear_reg() |> fit(score ~ age, data = evals.age)
model$fit
```

Call:


```
stats::lm(formula = score ~ age, data = data)
```

Coefficients:

| | |
|-------------|-----------|
| (Intercept) | age |
| 4.461932 | -0.005938 |

```
ggplot(evals.age, aes(x = age, y = score)) +
  geom_point() +
  labs(x = "Age", y = "Teaching Score",
       title = "Relationship between Teaching Score and Age") +
  geom_smooth(method = "lm", se = FALSE)
```


`geom_smooth()` using formula = 'y ~ x'




3. Assess the model assumptions from Task 2 by plotting the residuals against the explanatory variable and fitted values, respectively. Also, plot a histogram of the residuals to assess whether they are normally distributed.

```
evals.age <- evals |>
  select(score, age)
model <- linear_reg() |> fit(score ~ age, data = evals.age)
regression.points <- get_regression_points(model$fit)
ggplot(regression.points, aes(x = age, y = residual)) +
  geom_point() +
  labs(x = "Age", y = "Residual") +
  geom_hline(yintercept = 0, col = "blue", size = 1)
```


Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
i Please use `linewidth` instead.



```
ggplot(regression.points, aes(x = score_hat, y = residual)) +
  geom_point() +
  labs(x = "Fitted values", y = "Residual") +
  geom_hline(yintercept = 0, col = "blue", size = 1)
```



```
ggplot(regression.points, aes(x = residual)) +  
  geom_histogram(binwidth = 0.25, color = "white") +  
  labs(x = "Residual")
```



4. Perform the same analysis we did on life expectancy from the `gapminder` data set in 2007. However, subset the data for the year 1997. Are there any differences in the results across this 10 year period?

```
gapminder1997 <- gapminder |>
  filter(year == 1997) |>
  select(country, continent, lifeExp)

lifeExp.continent <- gapminder1997 |>
  summarize(median = median(lifeExp), mean = mean(lifeExp), .by=continent)
lifeExp.continent
```

```
# A tibble: 5 x 3
  continent median  mean
  <fct>      <dbl> <dbl>
1 Asia       70.3  68.0
2 Europe     76.1  75.5
3 Africa     52.8  53.6
4 Americas   72.1  71.2
5 Oceania    78.2  78.2
```

```
lifeExp.model <- linear_reg() |> fit(lifeExp ~ continent, data = gapminder1997)
lifeExp.model
```

parsnip model object

Call:

```
stats::lm(formula = lifeExp ~ continent, data = data)
```

Coefficients:

| | | | |
|------------------|-------------------|---------------|-----------------|
| (Intercept) | continentAmericas | continentAsia | continentEurope |
| 53.60 | 17.55 | 14.42 | 21.91 |
| continentOceania | | | |
| 24.59 | | | |

- Return to the `Credit` data set and fit a multiple regression model with `Balance` as the outcome variable, and `Income` and `Age` as the explanatory variables, respectively. Assess the assumptions of the multiple regression model.

```
Cred <- Credit |>
  select(Balance, Income, Age)
```

```
Cred |>
  skim()
```

Table 1: Data summary

| | |
|-----------------------------------|------|
| Name | Cred |
| Number of rows | 400 |
| Number of columns | 3 |
| Column type frequency: numeric | 3 |
| Group variables | None |

Variable type: numeric

| skim_variable | n_missing | complete_rate | mean | sd | p0 | p25 | p50 | p75 | p100 | hist |
|---------------|-----------|---------------|--------|--------|-------|-------|--------|--------|---------|------|
| Balance | 0 | 1 | 520.02 | 459.76 | 0.00 | 68.75 | 459.50 | 863.00 | 1999.00 | |
| Income | 0 | 1 | 45.22 | 35.24 | 10.35 | 21.01 | 33.12 | 57.47 | 186.63 | |
| Age | 0 | 1 | 55.67 | 17.25 | 23.00 | 41.75 | 56.00 | 70.00 | 98.00 | |


```
Cred |>
  cor()
```

```
      Balance      Income      Age
Balance 1.000000000 0.4636565 0.001835119
Income  0.463656457 1.0000000 0.175338403
Age      0.001835119 0.1753384 1.000000000
```

```
ggplot(Cred, aes(x = Age, y = Balance)) +
  geom_point() +
```

```
labs(x = "Age (in years)", y = "Credit card balance (in $)",
     title = "Relationship between balance and age") +
geom_smooth(method = "lm", se = FALSE)
```

`geom_smooth()` using formula = 'y ~ x'




```
Balance.model <- linear_reg() |> fit(Balance ~ Age + Income, data = Cred)
tidy(Balance.model)
```

A tibble: 3 x 5


| term | estimate | std.error | statistic | p.value |
|---------------|----------|-----------|-----------|----------|
| <chr> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 (Intercept) | 360. | 70.4 | 5.11 | 4.97e- 7 |
| 2 Age | -2.19 | 1.20 | -1.82 | 6.91e- 2 |
| 3 Income | 6.24 | 0.587 | 10.6 | 2.20e-23 |

```
regression.points <- get_regression_points(Balance.model$fit)
```

```
ggplot(regression.points, aes(x = Income, y = residual)) +
  geom_point() +
  labs(x = "Income (in $1000)", y = "Residual", title = "Residuals vs income") +
  geom_hline(yintercept = 0, col = "blue", size = 1)
```



```
ggplot(regression.points, aes(x = Age, y = residual)) +
  geom_point() +
  labs(x = "Age (in years)", y = "Residual", title = "Residuals vs age") +
  geom_hline(yintercept = 0, col = "blue", size = 1)
```



```
ggplot(regression.points, aes(x = residual)) +
  geom_histogram(color = "white") +
  labs(x = "Residual")
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.

